

Kuiyi (Cory) Gao

Chapel Hill, NC | kuiyigao@unc.edu
kuiyigao.github.io | github.com/KuiyiGao

Education

University of North Carolina at Chapel Hill

Jan 2026 – present

Undergraduate student; intended B.A. in Computer Science, expected May 2027.

- GPA: 3.68/4.00. Exchange student in Spring 2026; transferred in August 2026.
- Selected coursework: Natural Language Processing, Advanced Natural Language Processing, Machine Learning.

The Chinese University of Hong Kong

Sep 2023 – May 2026

Undergraduate study in Computer Science; transferred to UNC. Faculty of Engineering ELITE Stream.

- Selected coursework: Introduction to Deep Learning (graduate level), Artificial Intelligence, Data Structures, Formal Languages and Automata Theory.

Publications and Manuscripts

Wenxuan Wang, Xiaoyuan Liu, **Kuiyi Gao**, Jen-tse Huang, Youliang Yuan, Pinjia He, Shuai Wang, and Zhaopeng Tu.

Can't See the Forest for the Trees: Benchmarking Multimodal Safety Awareness for Multimodal LLMs. ACL 2025. Paper

Wenxuan Wang, **Kuiyi Gao**, Youliang Yuan, Jen-tse Huang, Qiuzhi Liu, Shuai Wang, Wenxiang Jiao, and Zhaopeng Tu.

Chain-of-Jailbreak Attack for Image Generation Models via Step by Step Editing. Findings of ACL 2025. Paper

Co-first-authored manuscript on **GUI-agent robustness under interface noise**. Under review, 2026.

Research Experience

Runtime Safety for Third-Party Agent Skills

Jun 2026 – present

MBZUAI UGRIP project and ongoing work; advisor: Dr. Steve Liu

- Built the runtime monitoring layer for third-party agent skills in a four-student UGRIP team, recording observed behavior and the evidence behind safety decisions.
- Developing typed Tickets that bind approved permissions, runtime observations, and decisions so an auditor can recompute a verdict. Manuscript in preparation.

GUI-Agent Robustness under Interface Noise

Sep 2025 – present

CUHK; advisor: Prof. Michael R. Lyu

- Co-first author of a benchmark studying how interface changes affect agents on web and Android tasks, including task completion, safety, and action trajectories.
- Designed evaluation measures for success, performance degradation, and clean–noisy action differences; contributed to the evaluation harness. Manuscript under review.

Research Experience (continued)

Kernel-Based Novelty Measures for Text

May – Aug 2025

CUHK summer research; advisor: Prof. Farzan Farnia

- Built a Python prototype to compare sentence-embedding distributions using RBF kernels and spectral entropy; examined centering and whitening for text novelty analysis.

Safety Evaluation of Image Generation Models

Jun 2024 – Feb 2025

Tencent AI Lab and CUHK; Chain-of-Jailbreak, Findings of ACL 2025

- Second author. Designed the study methodology, contributed benchmark data, and ran most of the safety and defense experiments.
- Wrote the manuscript and reviewer responses; evaluated how safety behavior changes across multi-turn image editing and tested a prompting-based defense.

Multimodal Safety Awareness

Dec 2023 – Oct 2024

Tencent AI Lab; MMSafeAware, ACL 2025

- Third author. Designed the evaluation methodology and built a benchmark of 1,500 image–prompt pairs across 29 safety scenarios, covering both unsafe and benign inputs.
- Evaluated multimodal models and ran vision-centric reasoning fine-tuning experiments within a study of nine models and three mitigation approaches.

Selected Projects

Long-Document Compression for Contract Analysis

2026

UNC COMP 586, Advanced Natural Language Processing

Code

- Implemented learned span saliency, document rendering, and an OCR-to-text question-answering pipeline for CUAD contracts. Prepared controls to compare span selection and rendering choices; evaluation is ongoing.

Visual Consistency in Imagist and Romantic Poetry

2025

CUHK ENGE3940, AI and Language

- Used repeated text-to-image generation and CLIP embedding distances to compare visual consistency across ten poems. Examined prompt-length control and the limits of treating model consistency as evidence about literary objectivity.

Movability Segmentation

2025

CUHK ELEG5491, Introduction to Deep Learning

Code

- Trained a DeepLabV3 model with a custom ASPP head on COCO-Stuff relabeled into four movability classes, exploring a visual prior for collision risk.

Honors and Skills

Honors: Professor Joseph Lau Scholarship of Innovation; CUHK College Entrance Scholarship; CUHK exchange scholarship.

Programming and tools: Python, C, Java; PyTorch, Hugging Face Transformers, Docker, Playwright.

Languages: Mandarin (native), English (TOEFL iBT 110, January 2025), Cantonese (intermediate).